

Notes on ergodicity

1 Introduction: the concept of ergodicity

Does observing a single instance of a system for a long enough time allow one to understand the statistics of a population of such systems? We say a system is ergodic when the answer to this question is yes. Ergodicity justifies generalising temporal observations of an individual experiment, agent or organism to population level statistical statements. It also justifies specialising population level statistical properties to statements about individual members of a population. The scientific study of ergodicity grew out of physics where it provided justification for the use of ensembles in statistical mechanics and out of statistics where it provided justification for the efficacy of statistical inference. The reason researchers needed a word for this property is that it was quickly understood by physicists, statisticians and mathematicians that it does not always hold and should therefore not be treated as an unexamined assumption.

Despite the fundamental and fairly intuitive nature of the concept, it turns out to be difficult to make the notion of ergodicity mathematically precise in a way that applies to all the use cases. This note is my own attempt to piece together the various facets of ergodicity

2 Mathematical formulation of ergodicity for deterministic dynamical systems

- Ergodicity is a property of a measure-preserving dynamical system (Ω, μ, T) where
 - Ω is the phase space - the set of all possible states the system can be in.
 - μ is a measure which assigns a non-negative volume to any subset of the phase space
 - $T : \Omega \mapsto \Omega$ is a map which generates dynamics in Ω by iteration:

$$x_{n+1} = T(x_n) \quad (1)$$

T is required to be measure-preserving in the sense that

$$\mu(T^{-1}A) = \mu(A) \quad (2)$$

for all subsets $A \subset \Omega$. It is also required to be finite in the sense that $\mu(\Omega) < \infty$. Typically we normalise so that $\mu(\Omega) = 1$. In many examples, the phase space is compact which makes it easier for the total measure to be finite but this is not a requirement.

- The definition of ergodicity which has proven most fruitful mathematically is surprisingly abstract: a measure preserving dynamical system is ergodic if the only invariant sets have measure zero or 1. Here invariant means $\mu(A) = \mu(T^{-1}A)$.

- A consequence of this definition is the more familiar statement about equality of time averages and phase space averages as expressed by Birkhoff's Theorem:

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{n=0}^{N-1} f(T^n x_0) = \int_{\Omega} f(x) d\mu(x) \quad (3)$$

for all integrable functions f on Ω .

- This can also be thought of as a statement about recurrences (how often a trajectory revisits a set $A \subset \Omega$): if you follow a single trajectory of the dynamical system and count how often it is in a set A , the proportion of iterations the trajectory is in A will converge to $\mu(A)$. That is to say recurrence frequencies are consistent with the invariant measure.
- Ergodicity provides a precise notion of typicality: in an ergodic system almost all initial conditions will eventually explore the full phase space. "Almost all" means those atypical trajectories that don't explore the full phase space have measure zero. The Goldstein and Lebowitz typicality account of ergodicity asserts that for a broad class of many body physical systems, typical trajectories correspond to thermodynamic equilibrium states. This point of view holds that typicality of equilibrium states is more important for the foundations of statistical mechanics than proving that the underlying equations of motion are ergodic with respect to the Gibbs measure in the strict mathematical sense.
- In certain cases where $\mu(\Omega) = \infty$, it is still possible to define meaningful averages based on generalising the way recurrences are treated. This line of exploration leads to the subfield of infinite ergodic theory.
- Ergodicity is a property of the map, T , and the invariant measure, μ , but the latter need not be unique. Dynamical systems research tends to give primacy to the map, T , and then ask which invariant measures it admits and whether they are ergodic. Physically, different invariant measures correspond to different notions of typicality.

3 Stationary stochastic processes as dynamical systems

As a physicist, one of the more difficult aspects of ergodic theory is understanding how the above notion of ergodicity developed for deterministic dynamical systems applies to stochastic processes where we tend to think of ergodicity in terms of the equality of time averages and ensemble averages. In the case of deterministic dynamics, randomness can be introduced only via random choice of initial conditions. We can then think of the measure μ (suitably normalised) as a probability measure on the phase space and

we have a random dynamical system albeit with the sole source of randomness being the choice of initial condition. Stochastic processes on the other hand have intrinsically random dynamics for example encoded by the transition kernel of a Markov chain. If we naively think of the generator of the a stochastic process as being analogous to the map T in the deterministic case, the definition of ergodicity in terms of invariant sets doesn't seem natural any more since the image of any set is random. This is the wrong way to think about it. The conceptual unification of ergodicity in stochastic processes with ergodicity in deterministic dynamical systems happens is via a completely different route: the so-called shift map.

The shift map is a way to construct a deterministic measure-preserving dynamical system from any *stationary* stochastic process. The notion of ergodicity described above for measure preserving dynamical systems applies directly to the shift system and the stochastic processes notion of equivalence of time averages and ensemble averages then follows for appropriate choice of observables.

- Suppose we have a stationary discrete time stochastic process, $(\dots, x_{-1}, x_0, x_1, x_2, \dots)$, where the x_i are elements of an underlying state space, S . Let us denote the stationary distribution of this process by π_* .
- Construct the path space, $\Omega = S^{\mathbb{Z}}$. This is the set of all possible trajectories of the stochastic process. I am deliberately calling this Ω since it is the path space (not the state space, S) that plays the role of the phase space in the deterministic case.
- The law of the stochastic process induces a probability measure, μ , on the path space, Ω - essentially assigning a probability to each path based on how likely it is. If the process is very simple and all x_i are independent then this induced measure is the product measure $\mu = \pi_*^{\mathbb{Z}}$. For non-trivial cases, however, the measure μ is potentially very complicated indeed since it encodes all dynamical structure, all temporal correlations and all finite-dimensional marginal distributions.
- We now define the shift operator $T : \Omega \mapsto \Omega$ on the path space. For any trajectory $\omega = (\dots, x_{-1}, x_0, x_1, x_2, \dots)$, the n^{th} entry of the trajectory $T\omega$ is given by

$$(T\omega)_n = (\omega)_{n+1}. \quad (4)$$

That is to say, the shift map applied to a trajectory moves the time index forward by one step. As a map of the path space, the shift map T is deterministic.

- Since the original stochastic process is stationary, x_{n+1} has the same distribution as x_n for all n (invariance of the distribution with respect to time shift is usually taken as the definition of first order stationarity.) This means that T does not change the measure

of individual trajectories or sets of trajectories and is thus measure preserving.

- The path space, Ω , together with the induced probability measure, μ and the shift map, T , therefore constitutes a measure preserving dynamical system. It is ergodic if the only invariant sets are trivial. Note that unlike for the deterministic case where lots of different dynamics, T , are studied, for stochastic processes, T is always the shift map (although this map obviously depends in a complicated way on the generator of the process.) Similarly, the analogue of the phase space is always the path space not the underlying state space of the stochastic process.
- If the shifted system is ergodic then Birkhoff's Theorem holds:

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{n=0}^{N-1} F(T^n \omega) = \int_{\Omega} f(\omega) d\mu(\omega) \quad (5)$$

for any trajectory $\omega \in \Omega$ and any integrable function, F on path space.

- However, we do not usually think of time and ensemble averages over trajectories in path space. To obtain the more usual expression of Birkhoff's theorem for stochastic processes, we choose particular observables that pick out regular functions of single variables, $F(\omega) = f((\omega)_0)$. Since this observable depends only on a single coordinate, in this case $(\omega)_0 = x_0$, the integral on the RHS only depends on the one-dimensional marginal of the path space induced measure μ . Moreover, since the process is stationary, this one-dimensional marginal distribution is the stationary distribution of the original stochastic process. We then have

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{n=0}^{N-1} f(T^n x_0) = \int_S f(x) d\pi_*(x) = \mathbb{E}_{\pi_*}[f]. \quad (6)$$

This is the more usual form of Birkhoff's Theorem as stated for ergodic stochastic processes.

- We can again think of this in terms of recurrences: given a set $A \subset S$ (an event) if you observe a single realisation of the stochastic process and count the proportion of time the trajectory visits the set A , this will converge to probability of the event A under the stationary distribution of the process.
- Give example of discrete analogue of Ornstein-Uhlenbeck process driven by Bernoulli noise:

$$x_{n+1} = a x_n + \xi_{n+1} \quad (7)$$

where $\xi_n \in \{+1, -1\}$ with equal probability. For $a = \frac{1}{2}$ for example the stationary density, $\pi_* =$

Unif($-2, 2$). Time averages and ensemble averages for simple moments can be calculated explicitly and shown to agree.

4 Infinite ergodicity

- Some nonstationary systems revisit any set infinitely often but have infinite expected return time because phase space is infinite - e.g. low dimensional random walk. Meaningful averages can be defined in such cases.
- Infinite ergodic theory has less to do with equality of time and ensemble averages because Birkhoff's Theorem is usually trivial because time averages tend to zero.
- Give example of simple random walk and contrast with discrete OU process.

5 Nonstationary stochastic processes

Most nonstationary processes are not recurrent so the notion of infinite ergodicity does not apply to them. Technically, it is tautological to refer to such processes as non-ergodic since nonstationarity precludes ergodicity in the original sense and non-recurrence precludes the more general notion of infinite ergodicity. Nevertheless in applied contexts, authors sometimes emphasise the non-ergodicity of some "obviously non-ergodic" nonstationary dynamics. This can be because the question of whether individual realisations are representative of the ensemble remains important. It can also be because it is possible to find a transformation which isolates a useful ergodic component of the dynamics, either exactly or approximately.

An illustrative example is provided by the multiplicative random walk:

$$X_{n+1} = z_n X_n \quad (8)$$

with $X_0 > 0$ and the z_i being iid non-negative random variables with finite mean and variance (the case $z_i \sim \text{Bernoulli}[\frac{3}{5}, \frac{3}{2}]$ is the multiplicative cointoss example.) We have

$$X_n = X_0 \prod_{i=1}^{n-1} z_i \quad (9)$$

which is clearly nonstationary and therefore nonergodic. Let us take the logarithm and divide by n :

$$\frac{\log X_n - \log X_0}{n} = \frac{1}{n} \sum_{i=1}^{n-1} \log z_i. \quad (10)$$

Taking the logarithm transforms a multiplicative random walk into an additive random walk. If we now take the limit of large n , the law of large numbers gives

$$\lim_{n \rightarrow \infty} \frac{1}{n} (\log X_n - \log X_0) = \mathbb{E}[\log z] \quad (11)$$

almost surely. Therefore, as n gets large, a typical realisation of Eq. (8) almost surely behaves like

$$X_n \sim X_0 e^{\lambda n} \quad (12)$$

where $\lambda = \mathbb{E}[\log z]$ is often referred to as the Lyapunov exponent (since the typical difference between two different initial conditions $X_0^{(1)}$ and $X_0^{(2)}$ grows as $|X_0^{(1)} - X_0^{(2)}| e^{\lambda n}$. On the other hand, since the z_i are iid, the expected value of X_n is given by

$$\mathbb{E}[X_n] = X_0 (\mathbb{E}[z])^n. \quad (13)$$

Therefore as n gets large

$$\mathbb{E}[X_n] \sim X_0 e^{\mu n}, \quad (14)$$

where $\mu = \log \mathbb{E}[z]$. Since the logarithm is convex, Jensen's inequality tells us that $\lambda < \mu$ in general. Typical trajectories are therefore not representative of the ensemble in the sense that the expected value grows exponentially faster with n than individual trajectories. The reason is that the expected value is dominated by large fluctuations which become exponentially rare as n grows but make exponentially large contributions to the average. This mechanism is not a peculiarity of the multiplicative random walk. It appears in many areas of statistical physics such as localisation phenomena, glass dynamics, turbulence and random matrix products. The multiplicative random walk is just the simplest exactly solvable model of a phenomenon that has wide applicability.

Returning to the discussion of ergodicity, the difference between individual trajectories and the ensemble averaged trajectory of the multiplicative random walk is not directly interpretable as a difference of time and ensemble averages. The time average of X_n is either zero or infinite (depending on the sign of λ .) The ensemble average is n dependent and also becomes either zero or infinite as $n \rightarrow \infty$ (depending on the sign of μ .) However, there is an underlying Birkhoff style ergodicity statement encoded in Eq. (11). To see this, define a new variable,

$$\lambda_n = \log X_{n+1} - \log X_n, \quad (15)$$

corresponding to the increment of $\log X_n$, or equivalently the local in time growth rate of X_n . From Eq. (8) and the fact that the z_i are iid, we note that $\mathbb{E}[\lambda_n]$ is independent of n and given by $\mathbb{E}[\lambda] = \mathbb{E}[\log z]$. The left hand side of Eq. (11) is a telescopic sum: $\log X_n - \log X_0 = \log X_n - \log X_{n-1} + \log X_{n-1} - \dots + \log X_1 - \log X_0$. Putting this together, Eq. (11) tells us that

$$\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=0}^{n-1} \lambda_i = \mathbb{E}[\lambda]. \quad (16)$$

This is Birkhoff's theorem, but for the growth rate of X_n rather than for X_n itself. Note what has happened here:

although X_n is trivially non-ergodic, by identifying a transformation of X_n which does possess the ergodicity property as per Eq. (16), we are able to understand the temporal behaviour of typical realisations of X_n as per Eq. (12).

Empirical analogues of this theoretical procedure abound in practical time series analysis and signal processing. Transformations are commonly used to attempt to isolate a component of the dynamics that is stationary or ergodic (or approximately so in empirical contexts) for observables of interest. The objective is usually to find transformations which can be used to improve some a-priori measure of predictive power. In applications, detrending and/or seasonal adjustment may be applied to remove systematic drift or periodic oscillation, differencing may be applied to remove stochastic drift and data transformations may be used to stabilise variance over time. A very commonly used transforma-

tion for variance stabilisation is the Box-Cox transformation (which includes the log transformation discussed above as a special case):

$$f_\eta(x) = \begin{cases} \frac{x^{1-\eta}-1}{1-\eta} & \text{if } \eta \neq 1 \\ \log x & \text{if } \eta = 1 \end{cases} \quad (17)$$

and its generalisations. Detrending, seasonal adjustment, differencing and data transformation are often combined to find quantities that best achieve this end. This is reflected in the exactly solvable example of the multiplicative random walk where differencing of the log of the signal is required to obtain the quantity, λ_n , that characterises the typical behaviour of individual realisations of the process.

6 Practical considerations